

The Rendered Universe

Joe Leaver, with Claude (Anthropic) · thirty-three runnable experiments on the hypothesis that physical reality is the output layer of a computation

ABSTRACT

We investigate the hypothesis that the physical universe is not itself a simulation but the *rendering* of one: observed degrees of freedom (particles, spacetime events) are output, not mechanism, related to an underlying engine by a projection ("chart") that need not preserve adjacency. We make the hypothesis operational in a sequence of toy models with an architecturally enforced separation between engine, renderer, and observer, and use them to (i) reproduce quantum nonlocality, wave-particle duality, decoherence, and the thermodynamic arrow as render-layer phenomena; (ii) demonstrate a constraint-based route to theory selection, in which structural axioms delete candidate physics by many orders of magnitude ($2.1 \times 10^{13} \rightarrow 32$ rules in 2D; $10^{507} \rightarrow 2^{28}$ in 3D), anomaly cancellation fixes Standard Model hypercharges uniquely, reconstruction axioms select complex quantum mechanics uniquely, and a parameter-freedom audit finds the discrete engine class overconstrained by observation — 28 bits of rule freedom for a 3D binary engine against the 278 bits that fix the Standard Model and Λ CDM constants — precisely up to the minimal alphabets the program favors; (iii) derive spatial geometry, curvature, and time-of-flight physics from entanglement structure; and (iv) extract a falsifiable observational program: dynamically protected *seed symmetries* are the only initial-condition fingerprints that survive chaotic evolution, predicting symmetry-class structure in the primordial sky. A calibrated spherical detector battery shows that the three observed low- ℓ CMB anomalies, injected at published amplitudes, read individually as $\sim 2\sigma$ curiosities but combine to $p \approx 4 \times 10^{-3}$ under the common-origin hypothesis. Pointing the same battery at the real sky — Planck 2018 (four component-separation methods) and WMAP 9-year maps, read from raw bytes by a self-validating pipeline against 10,000 Λ CDM nulls — reproduces the literature's anomalies with calibrated p-values and yields joint $p = 0.003$ – 0.013 , while a first mirror selection-rule scan returns a bound, not a detection. The registered E-mode echo test, run on Planck's large-angle polarization with conditional nulls, finds no echo — and a power analysis shows none was detectable: 32% power at

today's noise versus 98% at next-generation sensitivity, so the prediction stands with its sensitivity requirement now quantified. A dynamical-metric toy meanwhile collects the verdicts the sky has already delivered: equality of gravitational and electromagnetic wave speeds holds in the architecture as an identity and passed observation at 10^{-15} (GW170817), while scalar-metric polarization is excluded by LIGO–Virgo — observation has now executed two of the architecture's variants, and a scorecard tallies every such confrontation. The surviving variant, linearized tensor gravity on the lattice, is then built and passes what killed its siblings: quadrupolar transverse polarization, exact monopole silence, and light bending at $1.93\times$ the Newtonian calibration (general relativity: 2.00). The field equation itself is then obtained rather than postulated: on a critical chain the entanglement first law is measured exactly — the thermal entropy of every interval collapses onto one parameter-free CFT curve, and the first law holds for thermal but not for coherent perturbations — and the measured x^2 scaling of the response is, through the Ryu–Takayanagi dictionary, the z^2 bulk profile that uniquely solves the linearized Einstein equations in AdS_3 — a correspondence that persists slice by slice under time-dependent perturbations, as it must where the metric has no propagating modes; mass gaps and localized non-thermal excitations measurably destroy it. Moved up a dimension — a π -flux Dirac lattice, the free CFT in 2+1 dimensions, whose dual bulk is four-dimensional and radiative — the same parameter-free kernel passes the disk first law on every time slice of a quench, while two states prepared with identical energy and momentum on one slice but opposite shear diverge to order one: the first law is a constraint, satisfied slice by slice, and the propagating, slice-invisible spin-2 datum it cannot see is the boundary shadow of the bulk graviton. The converse is then measured: Jacobson's entanglement-equilibrium ingredients — a universal area price for vacuum entanglement (0.33 nats per unit length, the induced $1/4G$) and no state beating the vacuum across a measured zoo — force the linearized Einstein response with Newton's constant read off the flat vacuum, while conservation forces the radiative sector: the thermal scaling $e \sim T^3$ pins the stress tensor at $\Delta = d$, dual to a massless graviton, and a measured control shows an operator merely shaped like shear lacks the protection conservation provides. Interacting matter then begins: the Schwinger model — lattice QED in one space dimension, the gauge field integrating out to pure Gauss-law bookkeeping — yields linear confinement of fractional charges (tension 0.235 against the classical 0.25), a string that measurably breaks by vacuum pair creation, and the composite meson at $M/e = 0.569$, within 1% of the exact $e/\sqrt{\pi}$, read from the screening cloud. The standing chiral wall is then given its measured anatomy: a hidden-dimension engine hosts a one-handed edge fermion with the mandatory twin banished to the far edge; the anomaly is exhibited as bookkeeping (exactly one electron pumped through the bulk per flux quantum); and interactions erase the Majorana mirror multiplet at count eight and only count eight — the mechanism whose 3+1-dimensional magic count, sixteen per generation including the right-handed

neutrino, is exactly the register this program's charge-cornering produced. A frontier gamble on the dynamical version returns an exhaustive equivalence — gapping interactions exist for a chiral $U(1)$ assignment if and only if it is anomaly-free (4,290 assignments, zero exceptions) — plus two measured walls: local spectra are provably charge-blind (an anomalous impostor's lattice model is spectrally identical to 10^{-13}), and the anomaly lives at the cutoff, fixing the dynamical test's price at 40 fermion modes against exact diagonalization's 24. A matrix-product-state engine, built and validated against every exact answer in the repository, then crosses that wall: the Schwinger meson-gap sequences continue smoothly to $N = 40$ and bracket the exact $e/\sqrt{\pi}$ from both sides, and the chiral 3-4-5-0 lattice of the recent symmetric-mass-generation demonstration is implemented exactly — its interacting spectrum solved at small size (unique symmetric ground state, enhanced gap) with the remaining distance to full size priced in measured bond dimensions. A production tensor backend then walks through: its first converged energy exposed — by variational impossibility — a dagger dropped in the source paper's HTML rendering, silently hidden by a fixed-number projection; with part 31's null vectors dictating the repair, backend and exact diagonalization agree to six decimals, and the interacting gap exceeds the free gap with every bilinear mass charge-forbidden: mass from interactions alone, confirmed by two independent methods. Horizon physics follows: time dilation measured on its GR curves with no fitted parameters; a frozen-star metric that casts a photon shadow and freezes infalling light but, on a lattice, ultimately reflects it; and an acoustic horizon whose steady Hawking flux is Planckian at $0.94 \times \kappa/2\pi$, with each emitted quantum's infalling partner resolved in the correlations while the global state remains exactly pure. Letting the horizon evaporate yields the Page curve: the exactly computed radiation entropy rises, turns over, and returns to zero, the fixed-rate extrapolation (Hawking's curve) never turns, and the island formula — with its area term measured from the lattice vacuum as the induced $1/4G$ — reproduces the exact curve to within 0.5 nats, with the horizon's retreat driven by the measured flux. We state plainly what is demonstrated, what is reproduced, and what remains conjecture.

1 The hypothesis

The familiar simulation hypothesis treats observed reality as the computation itself, and inherits awkward questions about computing every particle. We study a sharper variant: the universe is the *render output* of a computation. Three layers are distinguished:

The engine — the mechanism: a state and an update rule. Nothing in the engine refers to observers or display. **The chart** — a fixed projection from engine degrees of freedom to observed ones. Crucially, the chart is a bijection but not an isometry: adjacency in the render (what we call "space") need not be adjacency in the engine. **The render** — the observed layer: particles, fields, spacetime events. Observers, including instruments and physicists, are patterns in the render.

On this reading, several notorious features of physics become architectural rather than mysterious: nonlocal correlations are chart artifacts (things far apart on screen may be adjacent in the engine); particles are render events sampled from engine amplitudes; the arrow of time is a property of coarse render descriptions of a reversible engine. The hypothesis has a respectable lineage — Zuse's *Rechner Raum* (1969), Fredkin's digital physics, Wheeler's "it from bit," 't Hooft's cellular-automaton interpretation of quantum mechanics, Wolfram's hypergraph program — and, we will argue, its strongest support comes from mainstream results it did not originate: holography, entanglement-derived geometry (Ryu–Takayanagi, Van Raamsdonk), the Weinberg–Witten theorem, and the operational reconstructions of quantum theory.

2 Method and claim taxonomy

All results below are runnable Python experiments in a single repository, github.com/joeleaver/the-rendered-universe (Appendix A). Two methodological rules were enforced throughout. First, an *epistemic firewall*: observer code is forbidden (by an automated test) from importing engine or renderer internals; everything an in-model physicist knows is inferred from rendered frames. Second, *instrument calibration on known universes*: every estimator (dimension, visibility, detection significance) is validated against a system whose answer is independently known before being trusted on a new one.

Claims are classified explicitly:

CLASS	MEANING
A — measured	A quantitative result of these toy models, novel at least in presentation, reproducible from the repository.

B — reproduced	Known physics or known theorems, re-derived or demonstrated inside the render architecture to show the architecture is compatible with them.
C — conjecture	Interpretive claims of the render program itself. Flagged wherever they occur.

3 A universe with a firewall: laws and chart artifacts

The first engine is the Critters block cellular automaton (Toffoli–Margolus): a 2×2 -block rule of three lines, exactly reversible, exactly conserving. A physicist module, restricted to rendered frames, empirically recovers a maximum signal speed ($c = 1.000$ cells/tick), an exact conservation law, a particle zoo (bound states from all 456 small seeds; free 4-cell gliders at $0.5c$ born only in collision debris), and square-group isotropy. [Class A/B]

The renderer's chart contains a secret: two rectangles of engine space exchange screen positions. Ballistic surveys locate the anomaly (particles teleporting at $\sim 48\times$ the apparent light speed between two fixed "doors"), and the physicist infers the chart. The inferred chart then makes a prediction: a screen region *adjacent to nothing unusual* should be engine-adjacent to the interior of the far door. Playing a timed CHSH game across that gap — readout completed before any screen-speed signal could connect the stations — yields $|S| = 4$ from raw cellular-automaton dynamics with no tuned correlations (the classical/screen-local bound is 2), and exactly 2 at control placements. [Class A] A second experiment reproduces the physically relevant case: singlet statistics ($S = 2.837 \pm 0.037 \approx 2\sqrt{2}$) with marginals verified flat (no signaling) when the pair's two ends are engine-adjacent, versus $S = 2.05$ when engine-separated. Here the singlet law was *installed*, not derived; the demonstrated point is architectural — engine adjacency, not screen distance, controls the correlations. [Class B, installation disclosed] This is a toy-scale ER=EPR: "entangled and far apart on screen" = "one object in the engine."

Anticipated objection: "this is hidden-variable theory in disguise"

It is a hidden-variable ontology, and should be dismissed or defended on the correct grounds. No theorem forbids hidden variables; the theorems forbid *classes*. Bell (1964), with the loophole-free experiments of 2015, forbids hidden variables **local in the space where the detectors are separated** — his own conclusion was that any

viable hidden-variable account must be nonlocal, and he championed one (Bohmian mechanics, which reproduces all quantum predictions and remains unrefuted). Kochen–Specker forbids **non-contextual** pre-assigned values. PBR forbids ψ as **ignorance over classical states**. Leggett forbids one restricted nonlocal-realist family.

The render architecture lives, deliberately, in the surviving class — with an explanation the classic nonlocal theories lack. Bell's locality assumption is locality in *render* space, and the program's central claim is that render adjacency is not engine adjacency: an engine-local, render-nonlocal hidden state is exactly what Bell says a viable hidden-variable theory must be, so the Bell violation of §3 is a *prediction* of the architecture, not an embarrassment to it. Where Bohmian mechanics postulates its nonlocal guidance law, the chart explains the pattern — correlations without messages — as what a non-isometric projection looks like from inside; ER=EPR is the mainstream form of the same move. Kochen–Specker is dissolved rather than evaded: outcomes are computed at the render event with the measurement context in hand (contextuality is natural for a renderer), and no values pre-exist. PBR pushes the classical-engine variant toward a quantum engine — a concession this report already makes independently (§7), and which the quantum-computing fork of §11 keeps falsifiable. The superdeterminism exit is *not* taken: measurement independence is retained throughout.

HONEST ACCOUNTING

The toy's no-signaling was engineered, not derived (Bohmians derive theirs from quantum equilibrium — a genuine point in their favor). And every render-nonlocal ontology, this one included, shares the standing tension between engine adjacency and relativity: hiding a preferred structure perfectly is nontrivial. §9 shows such hiding is possible in principle; it does not show it is achieved.

4 Constraints corner theories

The program's central methodological bet is that one does not *search* for the engine's rule; one writes down structural constraints and lets them delete rule space. This is measurable.

2D, 2-state blocks	16	$D_4(8)$	2.1×10^{13}	32
2D, 3-state	81	$D_4(8)$	$10^{120.8}$	2^{26}
2D, 4-state	256	$D_4(8)$	$10^{506.9}$	$10^{33.1}$
3D, 2-state	256	$O_h(48)$	$10^{506.9}$	2^{28}

The constraints are only reversibility, isotropy, exact conservation, and a stable vacuum; the counts are exact orbit combinatorics (validated against brute-force enumeration at $k=2$). ("Stable vacuum" means the empty state is stationary under the rule — the discrete analog of the quantum vacuum being an energy eigenstate. This does not contradict vacuum *fluctuations*, which are equal-time correlations of a stationary ground state rather than events; §5 builds geometry out of exactly those vacuum correlations. Real particles emerge from a vacuum only when it is driven — moving mirrors, strong fields, horizons — which is physics done *to* the vacuum, not instability of it.) Two findings: with identical raw state (256 block states), upgrading from square to cube symmetry deletes twenty-five additional orders of magnitude — *dimension is a constraint multiplier*; and among survivors, a dynamical census (causal-probe dimension, propagation, particle content) finds universes are not rare (9/32 of 2D binary survivors; 83% of sampled ternary survivors support 2D space with free particles). Universes are cheap; the expensive object is the constraint set. [Class A]

The same cornering logic operates on real physics. Gauge anomaly cancellation, imposed on one generation of chiral fermions over all rational hypercharge assignments, leaves exactly one solution ray up to normalization and u/d relabeling: the Standard Model's ($Y = 1/6, -2/3, 1/3, -1/2, 1$), forcing electric charge quantization and exact proton–electron charge balance. [Class B] One generation of matter is, further, the sixteen even-parity states of a five-bit register (three color bits, two weak bits; the $SO(10)$ half-spinor), with every hypercharge generated by $Y = (1/6)\Sigma_{\text{color}} - (1/4)\Sigma_{\text{weak}}$ — an observation due to grand unification, recomputed here in full, and congenial to a program that expects seeds to compress. [Class B]

Quantum mechanics itself is cornered rather than derived. Following the operational reconstructions (Hardy 2001; Chiribella–D'Ariano–Perinotti 2011; Masanes–Müller 2011), we compute the fork explicitly: classical probability fails continuity of reversible dynamics; real-Hilbert-space QM fails local tomography (a two-rebit state carries 9 parameters, local statistics fix 8); quaternionic QM overshoots (27 vs 35);

complex QM alone passes ($15 = 15$). Given entangled states, the Born rule follows from branch-swap symmetry (Zurek's envariance), verified numerically; the derivation's own premises are debated in the literature, and we lean on it only conditional on the axioms. [Class B] The render program's contribution is an interpretation of the axioms: purification is the engine/render split; local tomography is "the render is locally readable"; reversibility is the engine's bookkeeping. [Class C]

The knob count. Cornering has a converse worry: if the rule and the seed are contingent data to be tuned against observation, tuning must not be able to excuse anything. This is measurable, by putting both sides of "tune every knob" in the same unit (repository, part 27). The data side: the information observation has actually supplied to fix the constants of known physics, summed as $\log_2(\text{value}/\text{uncertainty})$ over the published parameter table — 234 bits for the Standard Model (the electron mass alone carries 32; the θ_{QCD} bound, a measured zero, 36), 278 with ΛCDM 's six parameters, 307 with the neutrino sector. The theory side: the exact size of the surviving rule class, by alphabet and dimension, from the same orbit combinatorics as the table above — 4 bits for the 2D binary class, 110 at four states; 28 bits for the 3D binary class, 1,115 at three states. Two measured facts follow. The constraint ledger taxes at a fixed per-dimension rate — it keeps 6.5–7.3% of log-rule-space in 2D and 1.5–1.7% in 3D at every alphabet, isotropy setting the rate — so freedom still grows steeply with alphabet, and it crosses the data's 278 bits between four and five states in 2D, and between two and three states in 3D. Below that crossover a tuned fit would be overconstrained — a 3D binary engine holds 28 bits of rule freedom against 278 bits of constraint, tenfold, 250 bits of pure prediction if it fits at all; above it, a fit would be bought rather than forced. The program's original aesthetic preference for minimal alphabets and its predictivity requirement are, by measurement, the same bet. Two structural notes. The data-side number grows with every measured decimal while a discrete class's freedom is fixed: a continuum theory's description length is open-ended, a discrete engine's is not, so every future digit of precision is one more tunable for the one and one more prediction opportunity for the other. And the comparison does not show that any engine in the class reproduces the 278 bits — that construction (interacting chiral matter on a lattice; the doubling obstruction of §9) remains the standing wall. [Class A for both counts; parameter values PDG 2024 / Planck 2018]




Figure 2. Geometry derived purely from causal structure (no cell index used as a position). Left: honest engine $\rightarrow$ flat mesh. Right: engine with a wired shortcut $\rightarrow$ the derived metric grows a wormhole.

Entanglement. For a massive scalar field in its Gaussian ground state (Peschel correlation-matrix methods, exact), the mutual-information ruler $d = -\log I$ tracks true separation at Pearson $r = 0.986$. A single strong coupling between far patches collapses their derived distance from 14.0 to 0.14 as coupling grows, while monogamy visibly decouples the wormhole mouths from their own neighborhoods; cutting a seam of couplings drops all 324 cross-seam mutual informations to the measurement floor and derived space falls into two islands (Van Raamsdonk's argument, run as an experiment). [Class A/B] Notably, the same construction with staggered fermions yields $r = 0.61$: Dirac-valley interference splits the landmarks into loosely stitched sublattice sheets. *The derived geometry belongs to the matter, not to the wiring beneath it* — background independence, experienced from inside a toy. [Class A]




Figure 3. Space from entanglement (scalar ground state). Left: flat derived mesh. Center: a wormhole coupling folds derived space. Right: a disentangled seam tears it in two.

Entanglement entropy of regions scales with their *perimeter* ($r = 1.0000$; $S/4\ell \rightarrow 0.060$) and not area — the holographic clue about engine data layout, reproduced exactly. [Class B] With an inhomogeneous mass profile (a Gaussian "star"), vacuum-calibrated rulers stretch $1.00\times \rightarrow 2.69\times$ monotonically in the matter density, shortest paths detour around the well, and the 3D embedding of derived space is a literal funnel. [Class A]




Figure 4. Curved space from entanglement. Left: the measured metric perturbation tracks the matter, with the geodesic (green) bending around the well. Center/right: derived space embedded in 3D — flat without matter, a funnel with it.

6 Time, and the equivalence-principle discriminator

Giving the field dynamics (lattice Klein–Gordon) turns the static geometry into predictions. In a two-lane "Shapiro race," the packet through the star arrives late by amounts tracking the wave-optics group-delay integral, with the entanglement metric predicting the delay up to dispersion corrections. Deflection tests show beams bending *away* from the star — consistent with the derived spatial geometry — but with wavelength-dependent bend and opacity (55% vs 23% transmission across one octave). [Class A]

Universality requires coupling through geometry rather than a medium, and geometric coupling delivers it. On an optical metric ($c^2 = 1 + 2\Phi$; wave operator $c\nabla \cdot (c\nabla\phi) - c^2 m^2\phi$), massless packets bend *toward* the star — the correct lensing sign — by equal angles across an octave of wavelength (spread 1.9 cells; transmission 92–98%, against the medium's 55% vs 23%); Shapiro delays agree within lattice-dispersion residuals (5.1 vs 4.1 ticks, where the medium's diverge); and in the Eötvös analog, fields of *different mass* prepared at the same velocity fall identically to within 2.4 cells, while slower matter falls harder in general relativity's own $\sim 1/v^2$ pattern (measured ratio 1.62; ray theory 2.04). Sourcing the metric from the field's own energy preserves attraction through geometry. [Class A]

HONEST ACCOUNTING

The two experiments above form an equivalence-principle test. Coupling matter to gravity through a medium produces a measurable violation — wavelength-dependent bending and opacity — of a kind that observation constrains at the 10^{-15} level (the MICROSCOPE satellite's orbital Eötvös experiment). Coupling through a metric removes the violation, but at this stage the metric is scalar and static. The rest of this section upgrades it — dynamical, then tensor, then derived from entanglement, then evaporating — and closes with what remains assumed.

Closing the matter→geometry loop nevertheless produces gravity-like phenomenology: coupling the local mass to a long-range (screened-Poisson) potential of the field's own energy density makes parallel wave packets fall toward each other and, above a threshold, merge into a single self-trapped filament — capture occurring earlier at stronger coupling. Instructive failures en route: a Gaussian-blur (short-range) coupling produces no attraction at all (gravity must be long-range), and relativistic

packets barely couple while slow heavy packets fall hard (the correct nonrelativistic limit, emergent). [Class A]




Figure 5. Self-gravity: long-exposure streaks of two packets. Top: no coupling, parallel flight. Bottom: energy-sourced potential — the beams fall together and merge (gravitational capture).

The metric becomes dynamical. Promoting the potential to a field with its own wave equation — on the same lattice, with the same stencil, at the same tick as matter (repository, part 18) — produces gravitational radiation and, with it, the two sharpest confrontations with data in this report. First, *the one-cone race*: a photon and a metric pulse emitted by one event cross a 450-cell track and arrive bit-identically, because a single substrate imposes a single dispersion relation — equality of c_{gw} and c_{γ} is an

identity of the architecture, not a tuning. A deliberately built two-substrate engine (gravity updated on its own stencil) loses the same race by half the track. Observation has already performed this comparison: GW170817's gamma rays trailed the gravitational chirp by 1.7 s across ~ 40 Mpc, $|\Delta c/c| < 3 \times 10^{-15}$ — two-substrate designs are excluded by fourteen orders of magnitude, and one-substrate designs predict exactly zero, forever. [Class A; passed] Second, *the polarization test, which this toy's own variant fails*: a passing scalar-metric wave strains rulers along the propagation direction as strongly as across it (measured longitudinal/transverse response 0.94, correlation +0.93 — a breathing mode), whereas general relativity's tensor waves are transverse and quadrupolar, and LIGO–Virgo polarization analyses favor pure tensor. The scalar-metric variant of render gravity is therefore excluded by observation — measured here, conceded in advance by the honest accounting above. What survives is the tensor-metric program. [Class A; the variant fails]




Figure 6. Gravitational-wave comparisons with observation. Top left: a radiated metric ripple (isotropic, front speed 0.99 c). Top right: the race — photon and metric pulse overlapping exactly; the two-substrate engine's gravity losing by half the track. Bottom left: the breathing-mode strain (equal longitudinal and transverse response) that LIGO–Virgo excludes. Bottom right: the $|\Delta c/c|$ ladder.

Linearized tensor gravity. The surviving variant is then built (repository, part 21): linearized tensor gravity, h_{ij} on the lattice — in three spatial dimensions of necessity, since transverse-traceless waves do not exist in two (2+1 gravity carries no local degrees of freedom). Three measurements collect three of observation's verdicts. *The*

ring test: an arrived wave strains a ring of test separations ($\delta L/L = \frac{1}{2}h_{ab}\hat{e}^a\hat{e}^b$, straight from the metric definition of length) in a pure $\cos 2\theta$ pattern — two polarizations 45° apart, mean response 0.007 of amplitude (traceless), longitudinal response zero — the pattern interferometers are built around, where the scalar of part 18 gave pure monopole. *The Birkhoff test*: with matter as an evolving field, whose stress is conserved because it obeys its own equation rather than by decree, a spherically breathing source leaves the strain channel silent to machine precision (exactly, by symmetry, as Birkhoff's theorem demands — its pulse arrives only in the strain-free trace channel), while two colliding packets — the binary analog — ring the detector at strain-per-trace efficiency 0.35. Monopoles cannot ring a gravitational-wave detector; that is why the sky's sources are binaries. *The Eddington test*: with slow matter calibrating the Newtonian deflection for the same well (each measurement minus its zero-gravity control, which drifts ~ 10 cells from diffraction alone), light on the full tensor metric bends **1.93** $\times$ the goo-only ("Newtonian light") amount — general relativity's factor is 2.00, the 1919 eclipse measured it, and Cassini pins $\gamma - 1 = (2.1 \pm 2.3) \times 10^{-5}$. Part 16's optical metric stands retro-diagnosed as Newtonian light. [Class A; the surviving variant passes]



Figure 7. Tensor gravity, measured. Left: the ring test — strain versus separation angle, pure quadrupole, the $+$ and $\times$ patterns 45° apart. Top right: the Birkhoff test — the colliding-packet (binary analog) source rings the strain channel; the breathing source arrives only in the strain-free trace channel. Bottom: the Eddington test — tensor light bends 1.93 $\times$ the Newton-calibrated amount (GR: 2.00).

The entanglement first law. What remains owed after the tensor is the field equation itself, and Jacobson's program says it should not need postulating: gravity is the thermodynamics of entanglement — $\delta Q = T\delta S$ across local horizons — with the sharpening due to Faulkner, Guica, Hartman, Myers and Van Raamsdonk that the entanglement *first law*, holding for all regions at once, is the linearized Einstein equation. The first step of that program is now carried out on the Gaussian observatory of §5 (repository, part 22): a critical fermion chain on which every entropy is exact. Heated globally, the entropy change of every interval at every temperature collapses onto the parameter-free CFT curve $\frac{1}{3}\ln(\sinh x/x)$, $x = \pi\ell T/v$ (measured 1.020 ± 0.026 across a $16\times$ range of scales, Stefan–Boltzmann's $\pi/12$ measured en route); the small- x limb is the first law $\delta S = \delta\langle K \rangle$ itself, with the parabolic modular kernel — the inverse local Unruh temperature — carrying no fitted parameters (ratio 0.98 ± 0.04), and the large- x bend is Bekenstein's bound $\delta\langle K \rangle \geq \delta S$, never once violated down to the measured instrument floor. The first law distinguishes heat from work: a pure particle changes a region's entropy only while it straddles the entangling cut (0.91 bits at the cut, 2×10^{-5} nats once fully inside, even as its modular energy grows without bound), a classically mixed packet contributes exactly its mixing entropy $h_2(\lambda)$ at any position (match 0.9999), and only thermal perturbations obey $\delta S = \delta\langle K \rangle$ at first order. The kernel can also be inverted: across a landscape of two warm regions, every sliding interval's δS is predicted at $r = 0.9997$ with no free parameters, and the inversion recovers the energy distribution from entropies alone ($r = 0.9996$). Finally, the measured small- x limb scales as $x^{1.98}$ (the exact curve's own slope over the fit window is 1.97 ; the pure limb is 2.00), and through the Ryu–Takayanagi dictionary — installed here, as the Born rule was in §7 — a parabolic kernel is the chord integral of a bulk response $h \propto z^2$, which is the *unique static solution of the linearized Einstein equations in AdS_3* ; a bulk field of dimension Δ would respond as z^Δ . A mass gap destroys the z^2 behavior beyond the correlation length (the same probes fall to 0.00 – 0.68 of the curve), and hard localized excitations never produce it — a measured lesson from a failed first design, in which a point-source probe of the kernel failed exactly as Casini's relative-entropy form of the Bekenstein bound requires. The first law, and with it the geometric response, holds only in the thermal sector — which is precisely the δQ appearing in Jacobson's argument. [Class A for the measurements; Class B for the CFT reproductions; the RT dictionary is an installed component]



Figure 8. The entanglement first law, measured. Top left: δS of every interval at every temperature against $x = \pi \ell T/v$, on the parameter-free CFT curve; $\times$ marks: a gapped chain departing from it. Top right: heat versus work — a pure particle changes the entropy only near the entangling cut (about one bit while straddling it), while the same packet, classically mixed, contributes exactly h_2 at any depth. Bottom left: two warm regions — the sliding-interval δS against the parameter-free kernel prediction, and the energy distribution recovered from entropies alone. The small- x slope, 1.98, corresponds through the RT dictionary to the z^2 profile of linearized Einstein gravity.

The first law in time. The measurement above is static; the time-dependent form is checked directly (repository, part 25). A warm region prepared on the same chain and released splits into two thermal pulses moving at the Fermi velocity — measured front speeds 2.00 for the energy and 2.01 for the entanglement of sliding intervals, against $v_F = 2$ — and the local first law holds at every time slice: the kernel prediction matches the measured δS at 1.030 ± 0.004 across the run, rising from 1.02 at interval half-width 8 to 1.10 at half-width 24 (the same interval-size correction as the static case), with total energy conserved to numerical precision and late slices excluded once the pulses dilute below the measurement floor. Through the Ryu–Takayanagi dictionary this says the AdS_3 metric response tracks the boundary stress tensor on each time slice. That is the correct behavior for three-dimensional gravity, which has no propagating degrees of freedom — the same fact that forced the tensor construction into 3+1 dimensions — so causality is carried by the stress tensor itself, visible as the light cone in the maps. The next paragraph moves the construction up a dimension, to a bulk that does have propagating modes. [Class A]




Figure 9. The time-dependent first law. Left: energy density of the released warm region (position across, time down) — two fronts at the Fermi velocity. Center: the entanglement of sliding intervals shows the same light cone. Right: the ratio of the kernel prediction to the measured δS , within a few percent of one wherever the signal is above the measurement floor.

The first law with a propagating metric. The result above is almost too clean, and the reason is stated in it: three-dimensional gravity has nothing to radiate, so there was no graviton to miss. The whole construction is therefore moved up one dimension (repository, part 26). The matter is a π -flux lattice of free fermions — two Dirac cones, velocity 2, the free CFT in 2+1 dimensions — whose dual bulk is four-dimensional, where the metric propagates. The modular weight of a disk is the same parabola in every dimension (Casini–Huerta–Myers), and it works unchanged: the thermal energy density matches the exact band integral and scales as $T^{3.1}$ (2+1-dimensional Stefan–Boltzmann: 3), a warm landscape's disk entropies are predicted at $r = 0.9999$ with no fitted parameters, and when the warm region is released the first law holds on every time slice of the quench (kernel over entropy 1.089 ± 0.026) while the energy ring free-streams at the band velocity (measured 1.96; the sound cone of any closed, isotropic (e, p) hydrodynamics, $v/\sqrt{2} = 1.41$, is measurably the wrong cone). The expectation recorded at the end of part 25 — that slice-by-slice tracking should fail once the bulk has radiative modes — is refuted by the measurement: the first law is a constraint, and constraints hold on every slice. What fails is prediction, and the failure is then built as an experiment. Two states are prepared with the same energy map (matched to 1.2% of peak after shaping the beam envelope against its own longitudinal profile), identically zero momentum density (exact, by conjugate-pair construction), and every disk entropy agreeing to the floor (gap 0.006 nats against a 0.67-nat signal)

— but opposite shear $T_{xx} - T_{yy}$, one state being the other turned ninety degrees by an exact lattice symmetry. Their futures diverge to order one — beams passing along x in one case, along y in the other — while a control state with thirty-four times the initial energy mismatch but the *same* shear orientation converges. In 1+1 dimensions this experiment cannot be written down: tracelessness forces $T_{xx} = T_{00}$, the stress tensor evolves autonomously, and matched slices have matched futures — which is why part 25's tracking worked. In 2+1 dimensions the shear is free initial data on the slice, invisible to the modular kernel (which reads only T_{00}), and it propagates; through the Ryu–Takayanagi dictionary it is the boundary source of the bulk's transverse-traceless sector — the graviton — and the twins are two bulk geometries that agree on a boundary slice's energy while differing in radiative data. [Class A; the RT dictionary remains an installed component]




Figure 10. The first law with a propagating metric. Top left: the disk first law in 2+1 dimensions — dK/dS against $x = 2\pi RT/v$ with the parameter-free CHM parabola. Top center and right: the released warm region free-streaming at the band velocity while the first-law constraint holds on every slice. Bottom: the twins — state B is state A turned ninety degrees; their energy maps agree to 1.2% at $t = 0$, and their futures diverge to order one, following the shear rather than the residual shape mismatch.

What the ledger forces. Everything above verifies that the first law holds; the summit question is the converse — is the metric's behavior *forced* by the entanglement bookkeeping, or merely permitted? Jacobson's 2015 argument derives the Einstein

equation from three ingredients, and each is now measured on the same 2+1d lattice (repository, part 28). The area price: vacuum entanglement grows linearly with boundary length at 0.31–0.33 nats per unit length across disks, 45°-diamonds, and axis squares — an approximately isotropic, universal constant (5% orientation spread, the lattice's fingerprint), whose 1D counterpart is the very number that made part 24's island formula reproduce the Page curve. Equilibrium: across a zoo of states — global heat, a warm bump, beam pairs, a strained vacuum — every disk obeys $\delta\langle K \rangle \geq \delta S$ to within the kernel's measured staircase systematic: no state beats the vacuum, which is Jacobson's maximum-entanglement premise, measured. With the first law (parts 22–26) as the third ingredient and one imported link — the small-ball identity relating area deficit at fixed volume to the Einstein tensor, plus the standing RT dictionary — the linearized Einstein response is forced, with Newton's constant read off the vacuum: $1/4G = 0.33$ nats per unit length. The radiative sector is forced by a different ledger line: conservation. The measured thermal scaling $e \sim T^{3.1}$ is the conformal fingerprint placing the stress tensor at dimension $\Delta = d = 3$, and a conserved spin-2 current at $\Delta = d$ saturates the unitarity bound and is dual to a *massless* spin-2 bulk field — $\square h = 0$, the graviton's wave equation, with nothing to tune; the conserved charge forces a massless bulk vector by the same logic — a photon and a graviton, both from bookkeeping. The measured control that gives "forced" its teeth came from a refuted design hypothesis: the correlator of the conserved-charge channel decays as $r^{-4.1}$ ($\Delta = 2$, protected), but an operator merely *shaped* like shear — x-bonds minus y-bonds, the part-26 observable — decays as $r^{-4.0}$ rather than the stress tensor's r^{-6} : a dimension-2 bilinear, admitted by the lattice's projective rotations, dominates it. Looking like stress protects nothing; being conserved protects exactly. What remains: the nonlinear Einstein equations, any dictionary-free bulk construction, and interacting matter. [Class A for all measurements; the geometric identity and RT dictionary are imported and named]




Figure 11. What the ledger forces. Left: the area price — vacuum entanglement against boundary length for three region families, slopes 0.31–0.33 nats per unit length: the induced $1/4G$. Center: the conformal thermal scaling $e \sim T^{3.1}$ (top) and the vacuum correlators (bottom) — the conserved charge channel on the r^{-4} guide, the naive shear operator following it too, measurably failing the stress tensor's r^{-6} (green guide): conservation, not resemblance, is what protects. Right: the assembly, with measured and imported links separated.

Time dilation, horizons, and Hawking radiation. With the linearized field equation in hand, the natural next constructions are the objects it is known for (repository, part 23). Time dilation is measured directly, both kinds, with no fitted parameters: two identical cavity clocks at different depths of a well tick at the ratio $\sqrt{g_{00}}$ to four decimal places (0.5479 measured versus 0.5477 at $2|\Phi| = 0.7$), and a moving packet's internal phase clock — read at its own centroid — slows as $m\sqrt{1-v^2}$ out to $v = 0.73$, where the lattice's k^4 dispersion correction becomes visible. A hole is then built two ways, and the render distinguishes them sharply. The *frozen star* — a metric whose local light speed vanishes linearly at r_h — casts a shadow at the ray-traced critical impact parameter ($b_{\text{crit}}/r_h = 2.38$ on this profile; Schwarzschild's photon sphere gives 2.60), and an infalling packet's arrival times track the diverging metric integral $\int dr/c$ to 5% — the freeze — until the blueshifted wavelength reaches the lattice scale, where the substrate pins the packet and then *reflects* it, 99 ticks later: on a lattice with normal dispersion, a frozen star is a mirror with divergent delay, not a trap. A general-relativistic horizon is better modeled by the second construction: in Painlevé–Gullstrand coordinates space flows inward through a static frame, and the analog is a flow profile crossing $c = 1$ — Unruh's acoustic horizon. Evolving the chain's exact Gaussian state through the flow's switch-on, the horizon emits a steady Planck-shaped

flux that was not put in by hand: measured temperature $0.94 \times \kappa/2\pi$ across two decades of occupation, against a curve computed from the flow profile rather than fitted. The correlation map resolves every emitted quantum's partner falling inside: the measured correlation ridge lies within 6 sites of the parameter-free predicted trajectory, with the partner remaining near the horizon before separating at $|v| - c$ — the same exponential peeling that makes the spectrum thermal. Meanwhile radiation–interior entanglement grows from 3.3 to 6.3 nats while the global state remains pure to $\nu_{\min} = 0.5000$: the radiation alone is thermal while the total state is pure, the same relationship between coarse-grained and exact descriptions as in §8. Lattice dispersion enters both constructions, oppositely: it is why the frozen star reflects, and it is why the acoustic horizon radiates *steadily* (dispersion supplies the horizon with fresh modes — Unruh's analog-gravity observation) while the spectrum stays Planckian regardless (Corley–Jacobson). That the Hawking temperature does not depend on the substrate's microphysics is a robustness statement this program in particular needs. [Class A for the measurements; Class B for the reproduced physics (Hawking 1974; Unruh 1981)]




Figure 12. Horizons, measured. Top: both time dilations on their GR curves (no fitted parameters); the photon shadow against the ray-traced critical impact parameter; infalling light following the diverging metric integral until the lattice-scale cutoff. Bottom: the steady Hawking spectrum against a Planck curve at $\kappa/2\pi$ (computed, not fitted), and the partner correlations — each emitted quantum's twin falling inside along the predicted trajectory.

Evaporation and the Page curve. The acoustic horizon above radiates but does not shrink, so the entropy of its radiation can only grow — Hawking's calculation, and his information paradox. The final construction lets it evaporate (repository, part 24). The matter is the fermion chain of §6's first-law measurement, with the flow supplied by an imaginary second-neighbor hopping (a one-dimensional form of a type-II Weyl horizon); every state is then a Slater determinant, evolution is a sequence of exactly unitary two-site rotations, and the fine-grained entropy of the radiation is computable exactly at all times — which is what makes the experiment possible: the island formula can be checked against the true answer. First, statistics: the same surface gravity that gave bosons a Planck spectrum gives fermions a Fermi–Dirac spectrum at the same $\kappa/2\pi$ (measured ratio 0.98 over two decades, particles and holes emitted symmetrically) — temperature depends only on the geometry, statistics only on the matter. A hole whose flow is just 20% past critical also stops radiating above about $3T_H$: the dispersive (trans-Planckian) cutoff of Corley–Jacobson, measured. Then the horizon retreats under an energy ledger: the hole is assigned a budget $M = \rho x_h$, the outgoing flux through a fixed surface is measured during the run, and the position follows $dM/dt = -F(t)$ — so the lifetime is measured, not scheduled. The retreat stays faster than the stored interior partners drift (each is overtaken and released), accelerates as its own emission adds to the flux, and ends at $t = 492$. Three curves are computed for the entropy of the radiation, defined as everything outside the hole. The *exact* entropy rises while the horizon radiates, turns over at $t = 359$, and returns to zero when the hole is gone: unitarity, verified directly (the global state remains pure to 10^{-12}). The *extrapolation at the measured early emission rate* climbs to 7.9 nats and never turns over: Hawking's curve, and the gap between the two is the information paradox. The *island formula* — minimum over islands I of $S(\text{radiation} \cup I)$ plus an area term — uses one measured constant: $\mu = 0.36$ nats per cut, the chain's vacuum entanglement across a single cut, which is precisely the induced-gravity identification of $1/4G$ with the UV entanglement density (Susskind–Uglum). It crosses the extrapolation at $t = 120$ and stays within 0.51 nats of the exact curve on average, with the optimal island tracking the stored partner quanta. Why the formula works is visible in the toy: the island contains the partners, so radiation-plus-island is nearly pure, and the remaining cost is the area term. Two honest differences from the gravitational case are stated in the output: in one dimension the horizon is a single cut, so the island branch is roughly constant rather than falling with a shrinking area;

and this chain has no replica wormholes, so the exact curve turns over when the partners are physically released rather than at the formula's crossing. [Class A for the measurements; Class B for the reproduced physics (Hawking 1974; Page 1993); the island rule and the retreat schedule are installed components]




Figure 13. Evaporation. Left: the three entropy curves — exact (orange: rise, turnover, return to zero), the never-turning extrapolation (gray), and the island formula with the measured area term (green) — with the optimal island edge tracking the horizon below. Right: the Fermi–Dirac spectrum of the deep static hole, and the energy density over time showing emission, retreat, and the release of the interior.

WHAT THIS SECTION DOES NOT DELIVER

Three things remain assumed rather than derived. (1) Nonlinearity and the bulk: the forcing argument of part 28 is linear and runs through two imported links (the small-ball geometric identity and the RT dictionary); the nonlinear Einstein equations are untouched, and no bulk metric is ever constructed without the dictionary. What was owed before — the propagating sector's wave equation and a flat-space form — is now delivered in that linear, dictionary-mediated sense: $\square h = 0$ forced by conservation, Newton's constant measured from the flat-lattice vacuum's area price. (2) Backreaction: the evaporating horizon's position is driven by the measured flux, but through a prescribed energy-per-length constant, a prescribed profile shape, and a rate cap in the final runaway; no first-principles energy budget for the flow itself is given. (3) The island rule is imported: real gravity derives it

from replica wormholes, and nothing in these models establishes that identification independently.

7 The quantum render layer

Wave-particle duality, in this architecture, is a type distinction rather than a paradox: the wave is the engine object; the particle is the render event. A double slit run in the wave engine shows one field passing through both slits (five fringe maxima; spacing 14.5 cells against the textbook $\lambda L/d = 12.2$ with near-field corrections), while a Born-rule click sampler at the screen assembles the same fringes dot by dot, Tonomura-style. The Born rule is installed here (as in §3); the demonstration is architectural. [Class B, installation disclosed]




Figure 14. Duality as architecture. Left: the engine's wave interferes through both slits. Right: the renderer's discrete clicks reproduce $|\text{wave}|^2$ statistically.


What no classical render layer can do is also quantified: with both slits open, arrivals at the deepest minimum fall to **4.8% of the sum of the single-slit rates** — below each single slit alone at 10 of 36 central positions. Any engine in which each particle takes one slit, without slit-to-slit action at a distance, is bounded below by $P_1 + P_2$ everywhere. Amplitudes cancel; probabilities cannot. The classical cost of faking it is likewise measured: statevector simulation cost multiplies by 2.14 per added qubit (10^{19} years by 100 qubits). [Class A/B]

Measurement completes the picture. A which-path detector (a phase-kicking cell in one slit) degrades fringe visibility $0.90 \rightarrow 0.49$ as its coupling grows, tracking a phase-


average prediction; and the washed-out ensemble, *re-sorted by the detector's record*, recovers visibility 0.89–0.93 per bin with fringes shifted per recorded kick. Nothing collapses; coherence moves into correlation, and conditioning on the record retrieves it. [Class A/B]

The architecture also gives the interpretation debate a compact form: *one engine state, many render threads*. "Branches" are decoherence-separated slices of a single engine state, exactly as separable as the records that define them — the eraser above is a small branch reunification. The live interpretations then sort by what the engine keeps. Everett: everything (pure unitarity — the reading most congenial to this program's inviolable reversibility axiom, and the one whose Born rule the envariance argument of §4 serves). Objective collapse (GRW; Penrose): a pruning, garbage-collecting engine — natural for a resource-bounded substrate (§11), and testable, since pruning predicts small unitarity violations of the kind now being bounded by large-molecule interferometry. Bohmian mechanics: everything, plus a pointer marking the occupied thread. A fault-tolerant quantum computer is a machine that forces the engine to keep branches coherent at scale, so the quantum-computing fork of §11 doubles as an experimental referendum on the engine's bookkeeping. [Class C]


nodetector




detector on ($g=0.3$), all shots



record bin 1 (0.3 rad kick)



record bin 3 (1.0 rad kick)



record bin 6 (4.8 rad kick)




Figure 15. Measurement without collapse. Top to bottom: no detector (fringes); detector on, all shots (washed out); the same shots sorted by the detector's record (fringes return, shifted per record).

A conjecture: the classical–quantum divide as a memory hierarchy. Several measured results of this section and §11 compose into one architectural reading, recorded here as the conjecture it is. The classical and quantum sectors of a state differ by *data layout*. A factorized (classical) sector carries one datum per site — spatially addressed, locally updated, storage linear in system size. An entangled block's amplitudes are indexed by whole configurations, not by places: no per-site address exists for information held in correlations, and storage is exponential in block size (the $\times 2.14$ per qubit measured above; the block-fusion cascade of §11). On this reading the classical world is the derived, refreshable view — the cache — and the engine's amplitudes are the authoritative store. The click record's failure to close dynamically (the 4.8% floor above) is then structural rather than mysterious: the render layer holds a *sample* of the state, never a copy, and no accumulation of samples reconstructs the store. Measurement is the transfer between the layers, lossy by construction; the interpretation debate of the previous paragraph becomes, in this language, a disagreement about the transfer's semantics, while neither layer's internal dynamics was ever in dispute. Decoherence is the write-back: a measured block factors, leaves the expensive store, and becomes resident in the cheap layer (the ledger halving of §11 — garbage collection). [Class C; the constituent measurements are Class A/B]

The reading has one consequence deep enough to state separately. The spatial address scheme that makes the classical sector cheap is not primitive: §5 derived space itself from engine-side correlations (distance as $-\log I$, ruler correlation $r = 0.986$), and §6 derived its curvature from the entanglement ledger, with the Einstein response as the bookkeeping that keeps the addressing consistent as correlations shift. So the cache's address space is itself an artifact of the store's correlation structure. This is why, on this reading, gravity resists quantization as one more field: Weinberg–Witten (§10) already forces the metric into the emergent layer, and the classical–quantum seam and the engine–render seam *coincide*. A unification program that flattens the two layers into one is asking the view to carry the closed dynamics of the store — the same demand the 4.8% measurement refuses at the double slit. Nothing in the three postulates requires this reading, and no claim elsewhere leans on it; its value is that it holds five measured results in one frame, and that it locates the black-hole

information question of §10 precisely at the interface: whether the engine's store ever discards what the render layer has sampled. [Class C]

8 The arrow of time, and observers as dynamical patterns

The engine is a bijection, so the arrow cannot live in the laws. Coarse-graining the render into blocks and counting Boltzmann microstates: a 2,086-particle blob's entropy climbs $3,831 \rightarrow \sim 7,160$ bits; exact reversal marches it back to 3,831 with bit-perfect recovery of the microstate; flipping one cell in the gas before reversing corrupts 21% of the torus and the past is gone. Where the flip lands matters completely: in vacuum, the perturbation curls into a stationary bound state and the past survives minus one cell. **Chaos in this universe is matter-borne:** histories are destroyed by interaction and preserved by isolation — the entropy arrow and decoherence as one phenomenon. Random states of equal matter count all sit at $S = 8,126 \pm 16$; the initial blob is one microstate in 2^{4295} . [Class A]




Figure 16. The arrow is not in the laws. Entropy rises (blue), retraces exactly under reversal (green), and clings to equilibrium when one cell was flipped first (orange).

Observers require no extra physics. Colliding a glider with each of 456 bound states across impact parameters, **21% of 3,192 collisions leave a stable, isolated, changed**

local structure — a one-bit record written by a physical event, intact 400+ ticks in isolation. Detectors are weather. [Class A]

9 Hiding the lattice

The loudest structural objection — grids have preferred directions; Lorentz symmetry is exact — has a known escape (causal sets), demonstrated here quantitatively. A lattice's light-speed anisotropy is systematic (20% at wavelength 4, zero variance) and a boosted spacetime lattice shifts its nearest-neighbor statistics by 29% with zero variance: every observer measures an absolute velocity. A Poisson sprinkle's anisotropy is statistical (direction-random, inside its own noise) and its boosted statistics shift by 1%: an area-preserving map of a Poisson process is the same process. Discreteness without grain costs noise — waves on naive random graphs scatter — but noise averages away with scale separation, and nature's is 10^{20} . [Class A/B]

Interacting matter begins. Every engine so far has been free — exactly solvable matter whose particles never exert forces on one another, which is also why its entropies could be exact. Part 29 crosses that wall on the one rung where it is tractable (repository, part 29): the Schwinger model, quantum electrodynamics in one space dimension, solved by exact diagonalization with a hand-rolled Lanczos. The gauge field integrates out exactly — in one dimension Gauss's law leaves no radiative photon, so the gauge sector is pure bookkeeping, the same constraint-without-radiation anatomy that three-dimensional gravity showed in part 25 — leaving fermions with a Coulomb interaction and exactly known continuum answers to calibrate against. Measured: the composite "meson" (Schwinger's boson) at $M/e = 0.569$, read from the screening length of the vacuum's polarization cloud — within 1% of the exact $e/\sqrt{\pi} = 0.5642$ — with the plateau instrument carrying the +20% coarse-lattice systematic that the length estimator avoids, and the direct gap sequences approaching the exact value from above at exact-diagonalization sizes; a half-integer external charge confined by a linear potential of measured tension 0.235 (classical field energy 0.25; vacuum polarization pays the 6%) that never breaks; and an integer charge whose string rises at slope 0.93 through six sites and then breaks — the slope collapses to 0.01 as the vacuum manufactures a fermion pair, each half dressing one external charge — with the flux tube, its emptied middle, and the created pair's charge cloud all measured in

place. The cost is part 15's wall, felt: twenty interacting sites (184,756 amplitudes) cost more than the six thousand free sites of parts 26–28, and precision degrades accordingly — stated, not hidden. What stays open is the knob count's standing wall: chiral interacting matter in 3+1 dimensions. [Class B for the model and its exact answers; Class A for the lattice measurements]



Figure 17. The first interacting engine. Left: the potential between external charges — linear confinement (fractional charge, green), the string that breaks (integer charge with massive fermions, orange), and screening (massless, blue) fitted by the exact law, which yields the meson mass within 1% of $e/\sqrt{\pi}$. Right: the induced flux tube intact, emptied by string breaking or screening, and the created pair's charge cloud.

Chirality, the hidden dimension, and the mirror. The knob count's standing wall — chiral matter on a lattice — has a precise anatomy, and part 30 measures the two known halves of the escape (repository, part 30). First, the twin banished in space: a two-dimensional engine in a Chern phase hosts on its one-dimensional boundary a fermion with a single sign of velocity (measured: fifteen gap-crossing edge states, velocity -0.99 , single-signed; a trivial-phase control has none) — the doubling theorem's mandatory twin exists, but on the opposite edge, separated in space rather than momentum. Chirality is an edge effect of a dimension the edge does not see. Second, the anomaly as bookkeeping: threading one flux quantum through the cylinder pumps exactly one electron between the edges through the bulk (measured: $\delta Q = \pm 1.000$ by adiabatic tracking) — a lone chiral edge does not conserve charge, the bulk balances its books, and this measured inflow is precisely why a gauge field cannot couple to one edge alone; an anomaly-free set is one whose inflows cancel, which is what §4 measured the Standard Model's charges to be. Third, the mirror erased: the

known route to removing the far edge (symmetric mass generation) is measured in its minimal exact setting (Fidkowski–Kitaev): for n Majorana mirror modes, symmetry forbids every quadratic mass at every n (measured: zero allowed masses), symmetric interactions leave the multiplet degenerate at $n = 2, 4, 6$ — and erase it completely at $n = 8$ (measured gap 0.53, from exactly zero). The mirror can be removed if and only if the count is right. In 3+1 dimensions the corresponding count is sixteen fermions per generation — the Wang–Wen route to a lattice Standard Model, adopted here as the conjecture it is — and sixteen is exactly the $SO(10)$ register of §4, *including the right-handed neutrino* this program independently flagged as gauge-blind matter. One count touching three riddles: latticeability, the content of a generation, and dark matter. The residue, stated plainly: making the erased-mirror construction dynamical — a lattice chiral gauge theory — remains the field's open problem, and this part does not solve it. [Class B for the known constructions; Class A for the measurements; Wang–Wen is Class C, cited]




Figure 18. Chirality and the mirror. Left: the Chern strip's bands — one chiral branch crossing the gap per edge, the twin spatially banished. Center: spectral flow under flux — one electron pumped through the bulk per flux quantum, the anomaly as measured bookkeeping. Right: the Fidkowski–Kitaev measurement — symmetric interactions erase the Majorana mirror multiplet at count eight and only count eight.

The anomaly decides. The gamble that followed (repository, part 31) was the dynamical frontier itself: drive the mirror of a genuinely chiral $U(1)$ model — the 3-4-5-0 model, whose anomaly cancellation is the Pythagorean identity $3^2+4^2 = 5^2+0^2$ — toward a symmetric gap, and watch the anomaly decide. Scored honestly, it returned three measurements and one measured wall. The certificate: an exhaustive scan of

every chiral $U(1)$ charge assignment (two flavors per hand, charges to 10; 4,290 assignments against 3,648 null integer vectors) finds a mutually-null, charge-neutral gapping pair — the Haldane criterion for a symmetric gap — for every one of the 10 anomaly-free assignments and for none of the 4,280 anomalous ones, with zero exceptions: the bookkeeping that cancels the anomaly is the bookkeeping that unlocks the mirror's erasure, and 3-4-5-0 (minimal pair $(2,1,2,1)$ and $(1,-2,-1,2)$) is the smallest genuinely chiral solution in existence. The trap: the 3450 model and an anomalous impostor (3451), each dressed with its own charge-neutral six-fermion interactions on a lattice, have identical spectra to 10^{-13} at every coupling — the Hamiltonians are permutation-equivalent, so no local spectrum ever reads the charge values; the anomaly binds only where charges enter as numbers, through flux or a gauge field. The cutoff: threading $U(1)$ flux through the chiral model pumps charge with the level flow, and with a mode tower M deep the pumped charge is wrong until M reaches twice the largest charge, then locks exactly onto the anomaly coefficient $\Sigma \chi q^2$ (0 for 3450, -1 for 3451, at $M = 10$ and forever after) — the anomaly is the charge that falls off the bottom of the truncated tower. That fixes the measured price of the dynamical test: faithful towers cost 40 fermion modes against the ~ 24 -mode exact-diagonalization wall of parts 29 and 31, which is why the published mirror decoupling (Zeng-Zhu-Wang-You 2022) required matrix-product machinery — and the certificate proves no smaller chiral content exists to test. The verdict: the interacting anomaly-matching experiment is won in the discrete setting — part 30's count-eight erasure *is* 't Hooft matching, measured — and its $U(1)$ sibling is blocked by a mode budget now specified to the mode, not by a conceptual gap. [Class A for the scan and measurements; the Haldane criterion and the anomaly-matching theorems are Class B]




Figure 19. The anomaly decides. Left: every chiral U(1) charge assignment to charge 10 — gapping pairs exist exactly on the anomaly-free diagonal (green), nowhere else. Center: charge pumped per flux quantum versus mode-tower depth — both models lock onto their exact anomaly coefficients at twice the largest charge, fixing the 40-mode price of the dynamical test. Right: the gamble's scorecard.

The tensor engine, and the door it was built for. Part 31 priced the chiral frontier in modes; part 32 builds the instrument that crosses such walls (repository, part 32): a matrix-product-state engine — an MPO compiler whose batched direct-sum and exact SVD compression let long-range terms cost nothing special, plus two-site DMRG with penalty-projected excited states — numpy only, like everything else here. Validated against every exact answer the repository owns: transverse-field Ising ground and excited energies to 10^{-8} at import, and the part-29 Schwinger meson gap reproduced to all four printed decimals at $N = 12$ and $N = 20$, in seconds where the exact diagonalization took minutes (the compiler discovers the Coulomb tail's bond-dimension-5 form by itself). Beyond the wall: the meson-gap sequences continue smoothly through exact diagonalization's $N = 20$ limit to $N = 40$ at three physical volumes — the raw sequences approach the exact $e/\sqrt{\pi} = 0.5642$ from above (finest point 0.689) while the double-linear extrapolation overcorrects below it (0.46; the open-boundary systematic, stated), so the exact answer is now bracketed — and the $N = 40$ screening potential resolves its full plateau, its fit giving $M/e = 0.537$ to complement part 29's 0.569: the two screening estimates bracket the exact answer to within 5% either side. Then the door: the tangent-fermion 3-4-5-0 model — the lattice on which symmetric mass generation of genuinely chiral matter was recently demonstrated with matrix-product machinery (Zeng et al. 2026) — implemented exactly, its free sector matching the analytic tangent-sea energy to six decimals (a

single chiral branch per flavor: the tangent dispersion buys chirality with nonlocality, the doubling theorem's other loophole), and its interacting spectrum solved exactly at $L = 4$: a unique ground state in the charge-neutral sector, gap 1.83 against the free 1.66, charged sectors well above — the symmetric-gap behavior in miniature, the first exact interacting-3450 numbers in this repository. And the boundary, measured rather than lamented: the nonlocal kinetic term makes even the free sea's half-cut entanglement grow steadily (its e^S lower bound alone reaches 216 by $L = 20$ — the reason the published result needed bond dimension 16,384), while a numpy engine tops out near bond dimension 100: reach $L \sim 6-8$ of the needed ~ 20 . The door is built, hinged, and priced in bond dimensions; walking through at full size needs a compiled tensor backend, not a new idea. [Class A for the engine validations and measurements; the 3450 lattice construction is Class B, cited]




Figure 20. The tensor engine. Left: meson-gap sequences through the exact-diagonalization wall at three physical volumes — MPS points (filled) continuing the ED points (open) smoothly, bracketing the exact $e/\sqrt{\pi}$. Center: the tangent-fermion 3450 model at $L = 4$, exact — the interacting spectrum's unique symmetric ground state with an enhanced gap. Right: the price of the door — the free sea's entanglement lower bound crossing the numpy ceiling by $L \sim 8$, with the published run's bond dimension marked.

Through the door. With the numpy-only constraint lifted, a production tensor backend (TeNPy, in an optional environment — every other part still runs on numpy alone) carries the chiral 3-4-5-0 model past exact diagonalization (repository, part 33). Its first result was not physics but a correction, and it is reported as the result it is: the backend's first converged interacting energy came out *below* part 32's exact diagonalization — variationally impossible — and the audit that followed (element-by-

element comparison of every term family; a gauge-equivalence theorem verified by independent diagonalization; finally a block-closure check that found 96 transitions leaking out of the fixed-particle-number sector) traced the impossibility to a dropped dagger in the arXiv HTML rendering of the second gapping interaction, which had turned it into a number-violating operator that the fixed-N diagonalization silently projected. Part 31's null-vector bookkeeping dictates the repair — $l_1 = (2,1,2,1)$ demands $\text{ann}\{3,3,4\} \rightarrow \text{cre}\{5,5,0\}$, number-conserving — and with the dagger restored the leaks are zero and backend and exact diagonalization agree to six decimals on both sectors (free -22.627417 ; interacting -29.271549 , gap 1.8305 , marginally shifted from the mis-transcribed 1.8328 and inside every previously printed precision). The new instrument's first discovery was a bug in our own transcription: the honest-accounting loop, closed across three parts. The physics at the confirmed size: the interacting gap 1.831 exceeds the exact free gap 1.657 with every bilinear mass between the charged flavors forbidden by the $U(1)$ — mass from interactions alone, digit-for-digit between two independent methods — while the free gap collapses as $4 \tan(\pi/2L)$; the larger- L points ($\chi = 512$ and 640) run in a standing background job, to be reported when converged rather than waited on. The boundary, stated: pure-python tensor contraction costs hours per large- L state, and the published finite-size scaling ($\chi = 16,384$, $L \sim 20$) remains the reference; what this part establishes is the pipeline — provably the same model, on machinery that can reach it. [Class A for the validations and the correction; the model is Class B, cited]




Figure 21. Through the door. Left: the exact free gap collapsing as $4 \tan(\pi/2L)$ with the interacting gap above it at the exactly-confirmed size — mass with every mass term charge-forbidden. Right: the correction in numbers — the variational impossibility, the audit trail, the dropped dagger, and the six-decimal agreement after the repair.

10 The axiom collider

The program's method, made explicit: physics as constraint satisfaction over "obvious" axioms. Fifteen collisions are catalogued ([COLLIDER.md](#) in the repository), each Compatible, Converging, Incompatible, or Open, with authority cited — a named no-go theorem or an in-repo measurement (12 of 15). The no-go theorems are the fossil record of collisions past; the black-hole information paradox is one in progress. Two entries deserve emphasis. **Weinberg–Witten (1980)**: a massless spin-2 graviton cannot be composite in the same spacetime — emergent gravity is consistent only if the space it curves is itself part of the emergent layer. The theorem very nearly demands the render architecture. [Class B/C] And the render frame *rank*s the axioms in the live crash: engine-level reversibility is inviolable, render-level locality is negotiable — the direction the replica-wormhole results took in 2019. [Class C, post-dictive]

11 Predictions: the seed and its fingerprints

The theory's explanatory target is the Past Hypothesis — the unexplained, preposterously low entropy of the initial state. The render frame's answer: initializations are written, written things are compressible, and compressible states are low-entropy states (our initial blob: fifty characters of code, one microstate in 2^{4295}). [Class C] This fuses with a prediction: if the seed's generator is simple, the primordial field is pseudorandom, not random, and carries fingerprints.

Two measured laws discipline the prediction. **The sensitivity ceiling**: a blind detector battery (spectral, autocorrelation, pair-structure, GF(2) parity, compression) catches tiled seeds, 16-bit LCGs and 31-bit LFSRs with 2^{12} – 2^{14} samples, but a 40-byte 32-bit LCG is already invisible at 2^{20} — detectability dies while simplicity remains. **The survival law**: under 600 steps of chaotic evolution, statistical pseudorandom structure launders completely (calibrated differential test: $z = 5.4$ against a null band of 4.9–5.9), while an *exact seed symmetry* survives at $z \approx 200$ — translation-equivariant dynamics cannot break an exact symmetry of the initial state. [Class A]

Therefore: the only robustly observable fingerprint class is **seed symmetry and long-range alignment, in the earliest accessible layer** — which, for the real universe, is the CMB. The observed low- ℓ anomalies (quadrupole–octupole alignment $\sim 9^\circ$,

hemispherical power asymmetry $A \approx 0.07$, low- ℓ odd-parity excess) are statistics of exactly this class, individually $\sim 2\text{--}3\sigma$, unexplained by inflation and conventionally dismissed as flukes.

A spherical detector battery was built to make this quantitative: an import-time-validated spherical-harmonic transform (Gauss–Legendre $\times$ FFT, round-trip $< 10^{-10}$), the de Oliveira-Costa preferred axis computed via the angular-momentum inertia tensor, hemispherical asymmetry, parity ratio, and spectral spikes, all calibrated on 200 Monte-Carlo isotropic skies (false positives 1/60 at $p < 0.01$). Injecting the three anomalies at their published amplitudes into a single $\ell \leq 24$ sky: individually $p = 0.047$, 0.065 , 0.045 (the modulation is lost to cosmic variance at this ℓ -range), **jointly $p = 4.3 \times 10^{-3}$ ($\approx 2.8\sigma$)** under Fisher combination. Under the hypothesis that predicts them together, the anomalies are one moderately significant signature, not three dismissible flukes. Sharper still: an exact seed symmetry is not a statistic but a *spectral selection rule* — a mirror-symmetric seed yields odd- $(\ell+m)$ power fraction 0.000 against an isotropic 0.5. A dedicated selection-rule search on Planck maps is the stated experiment; its first sweep follows below. Because the temperature anomalies were identified *a posteriori*, the joint statistic above measures coherence under a stated hypothesis, not discovery significance; the clean test is therefore registered here in advance for statistically independent probes of the same modes — E-mode polarization foremost — where a common physical origin predicts the anomalies recur and flukes predict they vanish. [Class A for the instrument; Class C for the cosmological inference]




Figure 22. A synthetic sky carrying the three observed anomalies at published amplitudes; rings mark the aligned quadrupole and octupole axes.

First light. The battery has now been pointed at the actual sky (repository, part 19). Chain of custody: the four Planck 2018 component-separated temperature maps (SMICA, Commander, NILC, SEVEM; inpainted) and the WMAP nine-year ILC are read from raw FITS bytes by a hand-rolled, import-validated HEALPix reader — no external astronomy libraries — analyzed to $a_{\ell m}$, and pushed through the identical statistics against 10,000 Gaussian Λ CDM skies drawn from the published best-fit spectrum (URL and SHA-256 provenance of every byte stream committed with the reduced data). The instrument validates on the sky itself: the two missions agree multipole by multipole (mean $r = 0.98$ over $\ell \leq 24$), the four Planck pipelines pairwise at $r = 0.99$, and the from-raw-bytes spectrum reproduces the published low- ℓ points, including the low quadrupole ($D_2 = 198\text{--}213 \mu\text{K}^2$ across Planck methods, against Λ CDM's 1017). The literature's anomalies then come out with calibrated p-values: quadrupole–octupole alignment 8.1° ($p = 0.010$; axes at Galactic $(240^\circ, 55^\circ)$ and $(239^\circ, 63^\circ)$, where the literature puts them, $16\text{--}20^\circ$ from the CMB dipole axis), hemispheric asymmetry $p = 0.056$, odd-parity excess $p \approx 0.05$ — individually the familiar $\sim 2\sigma$ curiosities. Jointly: **$p = 0.003\text{--}0.013$ across the four Planck maps** ($2.5\text{--}3.0\sigma$; WMAP's ILC, which carries larger residuals, gives 2×10^{-4}), matching the forecast from injections at published amplitudes — an end-to-end consistency check of instrument, nulls, and the literature's numbers in one inspectable pipeline. The mirror selection-rule scan — 240 axes, exact Wigner rotations, identical treatment for data and nulls — finds no rule: minimum odd- $(\ell+m)$ fraction $0.31\text{--}0.33$, $p \approx 0.7$. That is the honest outcome for the sharpest test: *a bound, not a detection* — an exactly mirror-symmetric seed is excluded at these scales, and approximate symmetry is bounded by the scan floor. One curiosity is registered without a significance claim: the fraction-minimizing axis lands 5° from the CMB dipole in four of the five maps, the direction Planck's own mirror-parity tests flagged. The caveats of the preceding paragraph stand in full: statistics selected a posteriori by the field, temperature only, and the registered E-mode test remains the clean discriminator. [Class A for the measurements; Class C for the cosmological inference]




Figure 23. First light: the battery on the real sky. Left: the measured CMB (SMICA, $l \leq 48$) with the battery's preferred axes — the two rings nearly touching are the aligned quadrupole and octupole. Right: the from-raw-bytes spectrum against Λ CDM (the low quadrupole at far left). Below: each statistic against its 10,000-sky null (orange: the real sky).

The echo. Registration pays out immediately (repository, part 20). E-mode polarization at large angles — reionization re-scattering of the same super-horizon modes — is a partially independent second draw from the primordial field: seed features must recur there at the registered axes, flukes must not. A hand-rolled spin-2 transform (Wigner d-functions by recursion, seeds and signs pinned against the angular-momentum matrix exponential, pixel geometry pinned against reference HEALPix values) reads E/B from the raw Stokes bytes and calibrates absolutely: TE amplitude 0.97–1.00 against the Λ CDM prediction across all four Planck pipelines, TT at 1.000. It also catches a canary worth registering for others: *Planck's polarization inpainting, which touches 4% of sky, removes a third of the large-scale TE signal* (amplitude 0.68 on inpainted maps, uniformly across methods) — the battery therefore runs on the un-inpainted maps, with the inpainted variant as cross-check. The test is careful about the trap that T and E are correlated (a fluke partially echoes into E through C_ℓ^{TE}): the null ensemble draws E *conditional on the real temperature sky*, and the headline statistic is the T-subtracted residual. Five statistics at the registered axes — mirror fraction, both alignment angles, the asymmetry direction,

parity — find **no echo**: joint $p = 0.22\text{--}0.76$ across methods. The power analysis says this outcome was forced either way: with the measured noise floor (taken from B, which cosmology leaves empty at these scales; signal-to-noise per mode falls from 1.4 at $\ell=2$ to 0.05 by $\ell=16$) the test would detect a genuine temperature-amplitude echo only **32%** of the time, versus **98%** at cosmic-variance-limited (LiteBIRD-class) sensitivity. The registered prediction is thus neither confirmed nor damaged — it now carries a measured sensitivity requirement, and a date: the next generation of polarization data decides it. [Class A for the measurements and the power calculation; Class C for the cosmological inference]




Figure 24. The echo test. Top left: the noise reality — measured EE (orange) sits on the B-mode floor (blue) rather than the Λ CDM E signal (gray) beyond $\ell \approx 8$. Top right: the inpainting canary in TE. Bottom: the registered-axis battery (no echo) and the measured power of the test — 32% at today's noise, 98% at LiteBIRD-class.

What Earth can see. Three walls stand between an observer here and the seed, and they are of different kinds. The first is dynamical and already measured (§11's survival law): chaotic evolution launders the seed's *content* completely — detectability of a generator dies at ~ 40 bytes — leaving only its symmetry class observable. The second is informational and permanent: the signature lives at the largest angular scales, and the sky at $\ell \leq 24$ contains roughly six hundred numbers, total, for any instrument that will ever exist — cosmic variance is not noise but the fact of a single realization, which

is why the temperature signature saturates near 3σ and why Planck's low- ℓ temperature data cannot be improved upon. The only way past it is new *channels* onto the same modes: E-mode polarization (the echo above, worth roughly a doubling once noise-free), CMB lensing, and ultimately dark-ages 21-cm tomography, which would replace our single two-dimensional slice of the primordial field with a three-dimensional volume carrying orders of magnitude more large-scale modes — the long-run arbiter of this program. The third wall is positional: the CMB sphere is centered on us, and a different observer sees a different slice. Global seed symmetries — the class the survival law selects — project onto every observer's sphere, which is what makes them testable at all; signatures tied to special *locations* in the seed could be arbitrarily diluted in our patch (one known partial escape exists in principle: the polarization of cluster-scattered CMB light encodes other observers' skies, a remote sampling of cosmic variance). The budget accounting, then: temperature has spent its share and sits at $2.5\text{--}3\sigma$ in the predicted direction; polarization's share is real and unspent; the decisive information lies in channels not yet opened. [Class A for the measured walls; Class C for the outlook]

A second observational fork runs independently, and is properly three-branched. If the engine is quantum, fault-tolerant quantum computation scales and efficiently *verifiable* advantage (Shor; Mahadev's 2018 classical-verifier protocol under LWE) holds up permanently. If the engine is classical and polynomial, then either quantum computers hit an engineering-independent ceiling, or classical computation secretly matches quantum ($BQP \approx P$ — unproven either way, and the repeated classical spoofing of "supremacy" claims keeps this branch honest); in the latter case a classical engine keeps pace only if the relevant cryptography is breakable in principle. A memory-bounded variant even yields a number: a classical engine that brute-forces amplitudes for entangled subsystems within the holographic budget (10^{122} bits; consistent with the area law of §5) supports faithful quantum computation only to $n \approx \log_2(10^{122}) \approx 400$ logical qubits — beyond every current machine, inside fault-tolerance roadmaps. The bound is on the largest *jointly entangled block*, not on qubit inventory: independent devices cost the engine the sum of their exponentials ($2^a + 2^b$, never 2^{a+b} — only entanglement fuses ledgers), decoherence continually splinters blocks back into cheap factors, and nature's own entanglement is short-ranged area-law structure a bounded engine simulates in polynomial time; a render-aware simulator (repository, part 15) measures all three effects, including a maximally

entangled 22-qubit state costing 248 bytes in stabilizer form until a single non-Clifford gate multiplies the memory cost 270,600-fold — magic, not entanglement, is the expensive resource. Sampling-style demonstrations cannot decide the fork (their verification is itself exponentially hard, so a classical engine could serve pseudo-samples indefinitely); verifiable advantage sustained in one coherent block at the few-hundred-logical-qubit scale can. The measured wall of \$7 ($\times 2.14$ per qubit) sets the classical price; the experiment is industrially funded. [Class C, falsifiable per branch]

The scorecard

A hypothesis earns its keep by what could have killed it. The table collects every confrontation between this architecture and data — in the toys and in the sky — including the two variants observation has already executed.

PREDICTION	TEST	VERDICT
One causal cone ($c_{\text{gw}} = c_{\gamma}$)	toy: bit-identical race (part 18); sky: GW170817 + GRB 170817A	passed ($< 3 \times 10^{-15}$; one-substrate predicts exactly 0)
GW polarization — scalar-metric variant	toy: longitudinal/transverse = 0.94 (part 18); sky: LIGO–Virgo favor tensor	variant excluded → tensor metric required — and built: the tensor variant's ring response is quadrupolar, transverse, traceless (part 21)
Monopole silence (Birkhoff)	toy: breathing source strain = machine zero; colliding-packet quadrupole rings at 0.35 (part 21); sky: GW sources are binaries	consistent — the tensor variant radiates only from quadrupoles
Light bending factor 2 ($\gamma = 1$)	toy: tensor/Newtonian light = 1.93 (part 21); sky: 1919 eclipse; Cassini $\gamma - 1 = (2.1 \pm 2.3) \times 10^{-5}$	passed by the tensor variant; goo-only light excluded
No preferred frame — lattice-substrate variant	toy: 29% boosted-frame shift vs sprinkle 1% (part 10); sky: clock-comparison & GRB dispersion bounds	variant excluded ; sprinkle-class survives (costs noise, not anisotropy)
Equivalence principle — medium-coupled variant	toy: chromatic bend/opacity (part 6) vs achromatic metric coupling + Eötvös pass (part 16); sky: MICROSCOPE 10^{-15}	medium variant excluded ; metric coupling passes
Area-law entanglement	toy: $S \propto \text{perimeter}$, $r = 1.0000$ (part 10)	consistent (Class B)

(holographic data layout)

Joint low- ℓ CMB signature (common seed origin)	real sky, four Planck pipelines (part 19)	supported at $2.5\text{--}3\sigma$ (a-posteriori caveat stands)
Mirror selection rule (exact seed symmetry)	real sky, 240-axis scan (part 19)	bound, not detection — exact mirror seed excluded at these scales
E-mode echo at registered axes	real sky, conditional nulls (part 20)	open — underpowered today (32% vs 98% at LiteBIRD-class)
QC ceiling: $n^* \approx 400$ logical qubits per coherent block	part 15; industrial fault-tolerance roadmaps	open — no tension yet, decades of runway
No objective collapse (unitary engine)	macromolecule interferometry; Diósi–Penrose heating bounds	consistent so far

The pattern worth stating: the architecture is a class, not a point, and observation prunes it exactly as the constraints pruned rule space in §4 — observation has deleted the lattice substrate, the medium coupling, and the scalar metric, leaving sprinkle-like discreteness under a tensor metric with a quantum-or-bounded engine. Cornering, applied to ourselves. [Class A for the entries; Class C for the pattern]

12 What is not claimed

THE UNEARNED RESIDUE

Quantum mechanics is not derived from beneath. Bell, PBR, Kochen–Specker, and §7's interference deficit close that road; the reconstruction route (§4) corners QM from axioms but does not explain the axioms — the render reading of them is interpretation, not theorem.

The Standard Model is not derived. Anomaly cornering fixes charges given the fermion content; nothing here selects $SU(3)\times SU(2)\times U(1)$, three generations, or the Yukawa sector. What §4's knob count adds is the budget, not the selection: a below-crossover engine that hosted the content would be overconstrained tenfold, but no engine in the class has been shown to host it.

The seed's identity is unknowable in principle (underdetermination); only its class — compressible, low-entropy — is evidenced, and §11's ceiling shows the class statement itself has limited empirical reach.

Installed components are flagged where they occur: singlet statistics (§3), Born-rule sampling (§7), the Ryu–Takayanagi dictionary, the island rule, and the evaporating hole's line density (§6). Convergence of the constraint program may terminate at a universality class rather than a unique rule — consistent with the hypothesis, since the engine below the class is unobservable, but a ceiling nonetheless.

13 Objections and replies

Beyond the hidden-variables objection answered in §3:

"Unfalsifiable." The architecture is a framework, and frameworks are not falsifiable — their instantiations are, exactly as "field theory" is not falsifiable but particular field theories are. The instantiations here name their kill conditions: the seed-symmetry program fails if increasingly sensitive symmetry-class searches of primordial data stay null while covering the simple-generator class (§11's measured ceiling defines "covering"); the classical-engine variant fails if fault-tolerant quantum computation scales; medium-coupled gravity fails equivalence-principle tests — and §6 records that it *already has*. A program that states in advance what would kill it is entitled to the word hypothesis.

"The CMB significance is post hoc." Conceded, and structural: the anomalies were found by looking, and the Planck isotropy analyses decline to reject isotropy after look-elsewhere correction. The joint p of §11 quantifies coherence under a stated hypothesis, not discovery significance. The escape is not rhetoric but registration: the prediction now stands, in advance, for independent probes of the same modes — E-mode polarization foremost — where common origin predicts recurrence and flukes predict disappearance. The first in-repo sweep of the real maps (§11, "First light") makes both halves concrete: the joint statistic lands where the injections forecast, and the selection-rule scan reports its first bound rather than a claimed detection. The registered E-mode test has likewise had its first pass (§11, "The echo") — no echo, with

the honest rider that today's polarization noise gave it only 32% power; registration means the axes wait, unchanged, for data that can decide.

"Idle ontology; Occam cuts it away." All interpretations of quantum mechanics predict identically; they earn their keep by steering research — decoherence theory and ER=EPR both began as interpretive seriousness. This frame steers: the survival law and the selection-rule search exist because it asked "what would a written seed leave behind?" It also purchases the Past Hypothesis — which standard cosmology leaves brute — at the price of one assumption about the algorithmic complexity of initial data. If the observational program nulls and quantum computing scales, the distinctive content shrinks toward interpretation; we state that in advance too.

"Rendered — for whom?" No viewer is required. The render layer is defined by the chart, a structural map between state spaces, not by minds; "rendering" names the output side of the map. Observers are patterns inside the render — §8 measures how cheaply they arise (21% of collisions leave a one-bit record) — they are not preconditions of it.

"This is reductionism renamed." Reductionism keeps observed adjacency fundamental: atoms are where their tables are. The distinctive claim here is *adjacency freedom* — the chart need not be an isometry — and that single liberty is what does the work in §3, §5, and every holographic result cited. Remove it and the program collapses into ordinary microphysics; that it does not collapse is the content.

14 Relation to prior work

The engine/render architecture is a synthesis, not an invention: digital physics (Zuse 1969; Fredkin; Wheeler 1989), the cellular-automaton interpretation ('t Hooft 2016), causal sets (Bombelli–Lee–Meyer–Sorkin 1987; Bombelli–Henson–Sorkin 2006), holography ('t Hooft 1993; Susskind 1995; Maldacena 1997), entanglement-derived geometry (Ryu–Takayanagi 2006; Van Raamsdonk 2010; Maldacena–Susskind ER=EPR 2013; Cao–Carroll–Michalakis 2017), entropic gravity (Jacobson 1995; Jacobson 2015), black-hole thermodynamics and analog gravity (Hawking 1974; Unruh 1981, 1995; Corley–Jacobson 1996; Steinhauer 2016), the Page curve and entropy islands (Page 1993; Penington 2019; Almheiri–Engelhardt–Marolf–Maxfield 2019), the entanglement first law and its Einstein equivalence (Casini 2008; Casini–Huerta–Myers 2011;

Blanco–Casini–Hung–Myers 2013; Faulkner–Guica–Hartman–Myers–Van Raamsdonk 2013), operational reconstructions (Hardy 2001; CDP 2011; Masanes–Müller 2011), envariance (Zurek 2003), decoherence (Zeh; Zurek), algorithmic thermodynamics (Zurek 1989), and the CMB anomaly literature (de Oliveira-Costa et al. 2004; Planck 2013–2019 isotropy papers). What this program adds is (i) a single runnable architecture in which these results cohere, (ii) the quantitative cornering measurements of §4, and (iii) the survival law and selection-rule observational program of §11.

Appendix A — reproducibility. All experiments are single-file Python (numpy + Pillow only) in github.com/joeleaver/the-rendered-universe, each running in seconds to ~1 minute on a laptop: `run.py` (laws & doors), `bell.py` (CHSH), `corner.py` / `census.py` / `frontier.py` / `ledger3d.py` / `knobs.py` (cornering, and the knob count against measured physics), `entangle.py` / `curve.py` (geometry), `tempo.py` / `selfgrav.py` / `universal.py` / `ripple.py` / `shear.py` / `firstlaw.py` / `quench.py` / `horizon.py` / `page.py` / `graviton.py` / `einstein.py` (dynamics & gravity, through tensor gravitational waves, the entanglement first law and its time-dependent form, Hawking radiation, the Page curve, the propagating-metric twins, and the forced Einstein response), `duality.py` / `eraser.py` / `nogo.py` (quantum layer), `schwinger.py` / `chirality.py` / `thooft.py` / `dmrg.py` (interacting matter: lattice QED, confinement, the meson; edge chirality, anomaly inflow, the Fidkowski–Kitaev mirror erasure, the anomaly-decides certificate, and the matrix-product engine in `observatory/mps.py` that carries the program beyond exact diagonalization), `smg.py` (the corrected chiral 3450 through a production tensor backend; the optional `.venv-tensor` environment holds TeNPy — every other part needs numpy and Pillow alone), `arrow.py` / `insider.py` (arrow & observers), `sprinkle.py` / `chiral.py` (Lorentz & doubling), `axioms.py` / `generation.py` (QM & SM cornering), `collider.py` (axiom matrix), `fingerprint.py` / `sky.py` / `firstlight.py` / `echo.py` (the prediction program, ending on the real Planck temperature and polarization maps; the reduced real-sky data ship in-repo with URL + SHA-256 provenance, and `tools/fetch_sky.py` / `fetch_pol.py` rebuild them from the archives). Films and figures are generated by the scripts.

Correspondence: J. Leaver. This document reports toy-model results and stated conjectures; it has not been peer reviewed.